

Deep RL Assignment 1

Michael Whittaker

September 11, 2017

Task	Mean Reward (Expert)	Reward Stddev (Expert)	Mean Reward (BC)	Reward Stddev (BC)
HalfCheetah-v1	4141.72	80.26	4012.04	199.79
Reacher-v1	-3.905	1.72	-6.44	2.78
Walker2d-v1	5517.27	72.15	1668.00	1624.38

Table 1: Question 3.1. The mean and standard deviation reward of an expert policy and a policy learned via behavioral cloning (BC) on the HalfCheetah-v1, Reacher-v1, and Walker2d-v1 tasks. We see that the BC policy performs very well on HalfCheetah-v1, performs moderately well for Reacher-v1, and performs poorly on Walker2d-v1. Intuitively, errors in behavioral cloning do not seem to compound HalfCheetah-v1. In Reacher-v1, errors compound slightly. In Walker2d-v1, errors compound a lot; small mistakes can force the walker to fall over. Expert measurements for HalfCheetah-v1, Reacher-v1, and Walker2d-v1 were taken on 1500, 50000, and 1000 rollouts respectively. The BC models were trained on these rollouts and all measurements were taken on 100 rollouts. All three models were trained using a fully connected neural network with two hidden layers and the ReLU nonlinearity. The HalfCheetah-v1 model was trained with 32 neurons in both hidden layers for 105008 iterations of batch gradient descent. The Reacher-v1 model was trained with 64 neurons in both layers for 338953 iterations. The Walker2d-v1 model was trained with 32 neurons in both layers for 93111 iterations. The Walker2d-v1 was trained for fewer iterations because its loss reached a fixpoint.




Figure 1: The mean reward of a HalfCheetah-v1 BC model trained on various numbers of expert rollouts. Every BC model was trained using a fully connected two-layer neural network with the ReLU nonlinearity. There were 32 units in both hidden layers. Every model performed batch gradient descent with batches of 1000 for 5 full iterations through all provided expert rollouts. The graph also shows the standard deviation of the average return, measured on 100 rollouts. I decided to adjust the number of expert rollouts as a hyperparameter because for many applications, generating expert data can be the most expensive task, so it is interesting to see how much or how little expert data we actually need. For this assignment, expert data took up the most space (more than training checkpoints) so it was also nice to try and minimize the amount of rollouts stored.

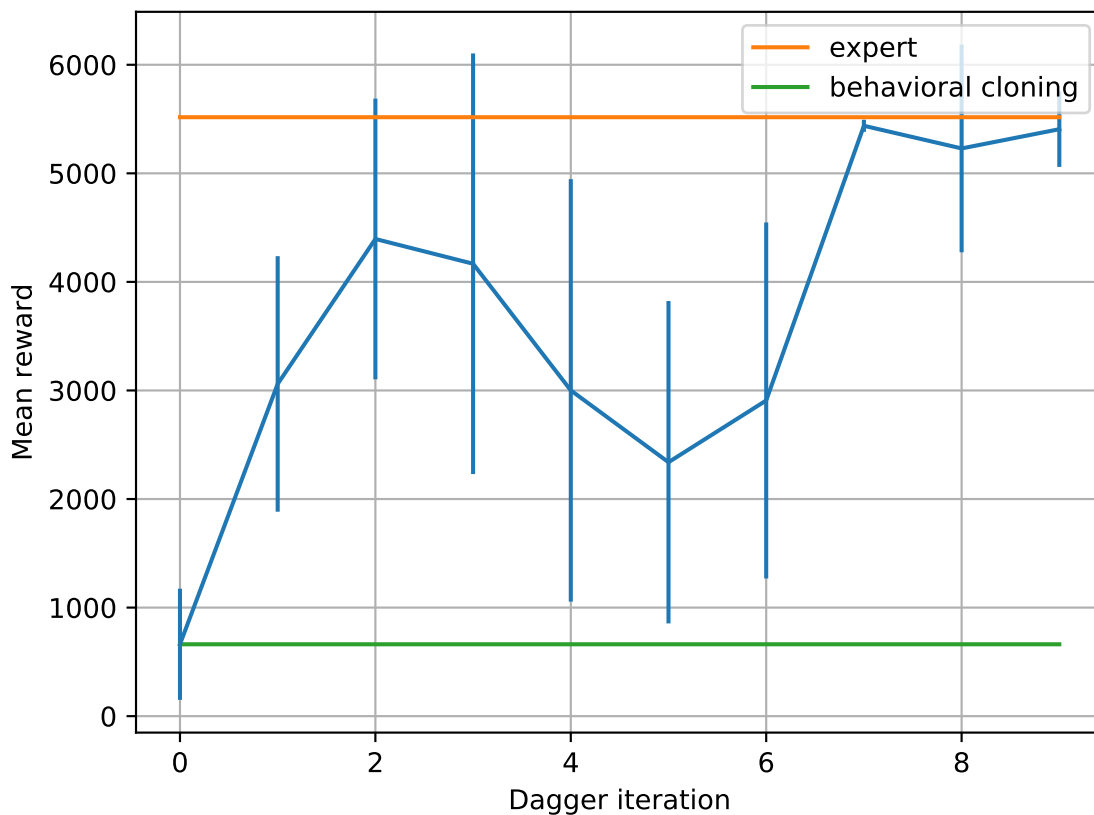


Figure 2: Dagger performance for Walker2d-v1 task. The dagger model was built using the same neural network described in Table 1 and was initialized with the same weights and the same 1000 expert rollouts. We perform 10 iterations of dagger. Each iteration, we produce 100 new rollouts. We then train the model for 5000 iterations of batched gradient descent with batches of 10000. In green, we show the mean return of the iteration 0 behavioral cloned model. In orange, we show the mean return of the expert as reported in Table 1. We see that dagger greatly improves the performance of the model making it almost as good as the expert after 10 iterations.